

Course: MMDS

Lab 01

Set up Hadoop


Fill answers of the questions below in the given tables.

Your screenshots must contain commands for required operations.

Question 1:

Set up Hadoop following this [tutorial](#) and then take a screenshot of the hdfs-site ([localhost:9870](#)) which shows information for NameNode Journal Status, NameNode Storage, DFS Storage Types (Overview page).

For example,



The screenshot shows the Hadoop HDFS site overview page. The browser address bar indicates the URL is localhost:9870/dfshealth.html#tab-overview. The page content is as follows:

NameNode Journal Status

Current transaction ID: 371

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-ntan/dfs/name)	EditLogOutputStream(/tmp/hadoop-ntan/dfs/name/current/edits_inprogress_00000000000000000371)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-ntan/dfs/name	IMAGE_AND_EDITS	Failed

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	19.56 GB	346.14 MB (1.73%)	9.09 GB (46.49%)	346.14 MB	1

My screenshot goes here

Namenode information

localhost:9870/dfshealth.html#tab-overview

NameNode Journal Status

Current transaction ID: 3

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-hoang/dfs/name)	EditLogFileOutputStream(/tmp/hadoop-hoang/dfs/name/current/edits_inprogress_0000000000000000003)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-hoang/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	24.48 GB	24 KB (0%)	14.92 GB (60.94%)	24 KB	1

Hadoop, 2019.

Question 2:

Set up Hadoop following this [tutorial](#) and then take a screenshot of the yarn-site ([localhost:8088](#)) which shows information for nodes ([http://localhost:8088/cluster/nodes](#)). For example,

localhost:8088/cluster/nodes

Logged in as: dr.who

Nodes of the cluster

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	8 GB	0 B	0	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries

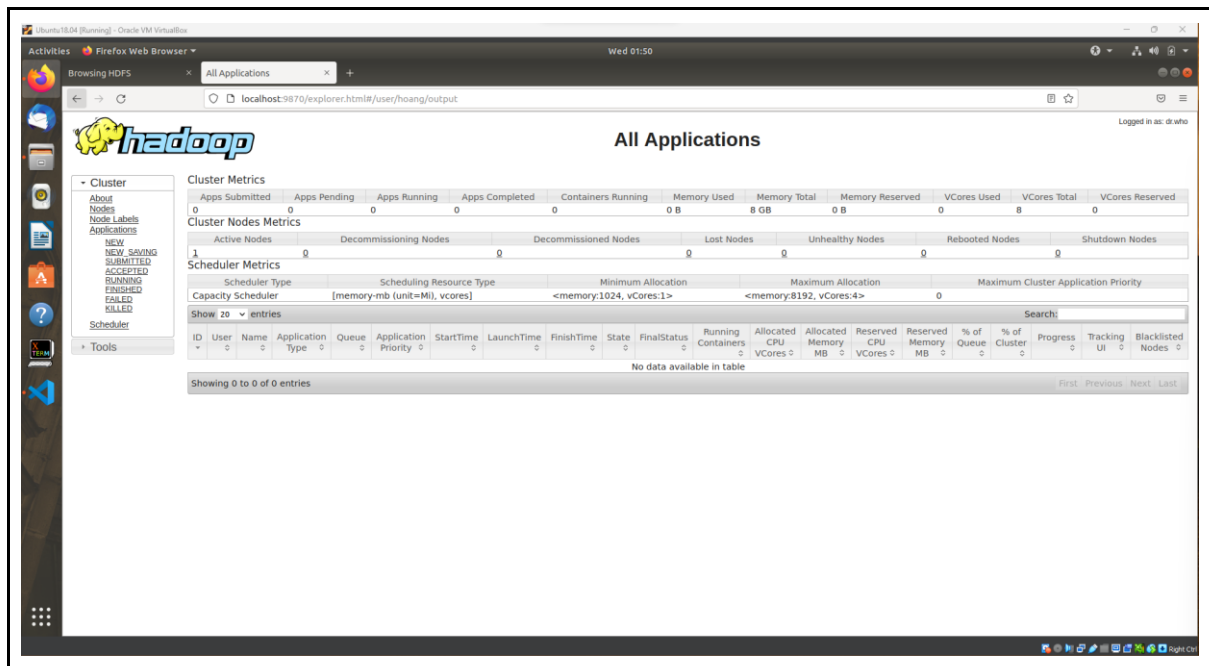
Search:

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Allocation Tags	Mem Used	Mem Avail	VCores Used	VCores Avail	Version
/default-rack	RUNNING	ubuntu:34741	ubuntu:8042	Tue May 11 12:30:58 +0700 2021	0	0 B	8 GB	0	8	3.2.1			

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

My screenshot goes here



Submission Notice

- Export your answer file as pdf
- Rename the pdf following the format:
lab01_<student number>_<Full Name>.pdf
E.g: **lab01_123456_NguyenThanhAn.pdf**
If you have not been assigned a student number yet, then use 123456 instead.
- Careless mistakes in filename, format, question order, etc. are not accepted (0 pts).